

Chapter 4 The Tidyverse

4.7 Summarizing data

An important part of exploratory data analysis is summarizing data. The average and standard deviation are two examples of widely used summary statistics. More informative summaries can often be achieved by first splitting data into groups. In this section, we cover two new dplyr verbs that make these computations easier: summarize and group_by. We learn to access resulting values using the pull function.

4.7.1 summarize

The summarize function in dplyr provides a way to compute summary statistics with intuitive and readable code. We start with a simple example based on heights. The heights dataset includes heights and sex reported by students in an in-class survey.

```
library(dplyr)

## 
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
## 
##     filter, lag

## The following objects are masked from 'package:base':
## 
##     intersect, setdiff, setequal, union

library(dslabs)
data(heights)
```

The following code computes the average and standard deviation for females:

```
s <- heights %>%
  filter(sex == "Female") %>%
  summarize(average = mean(height), standard_deviation = sd(height))
pander::pander(s)
```

average	standard_deviation
64.94	3.761

This takes our original data table as input, filters it to keep only females, and then produces a new summarized table with just the average and the standard deviation of heights. We get to choose the names of the columns

of the resulting table. For example, above we decided to use average and standard_deviations, but we could have used other names just the same.

Because the resulting table stored in s is a data frame, we can access the components with the accessor \$:

```
s$average
```

```
## [1] 64.93942
```

```
s$standard_deviations
```

```
## [1] 3.760656
```

As with most other dplyr functions, summarize is aware of the variable names and we can use them directly. So when inside the call to the summarize function we write mean(height), the function is accessing the column with the name “height” and then computing the average of the resulting numeric vector. We can compute any other summary that operates on vectors and returns a single value.

For another example of how we can use the summarize function, let’s compute the average murder rate for the United States. Remember our data table includes total murders and population size for each state and we have already used dplyr to add a murder rate column:

```
murders <- murders %>% mutate(rate = total/population*100000)
```

Remember that the US murder rate is not the average of the state murder rates:

```
pander::pander(summarize(murders, mean(rate)))
```

mean(rate)
2.779

This is because in the computation above the small states are given the same weight as the large ones. The US murder rate is the total number of murders in the US divided by the total US population. So the correct computation is:

```
us_murder_rate <- murders %>%
  summarize(rate = sum(total) / sum(population) * 100000)
pander::pander(us_murder_rate)
```

rate
3.035

This computation counts larger states proportionally to their size which results in a larger value.

4.7.2 Multiple summaries

Suppose we want three summaries from the same variable such as the median, minimum, and maximum heights. We can use summarize like this:

But we can obtain these three values with just one line using the quantile function: `quantile(x, c(0.5, 0, 1))` returns the median (50th percentile), the min (0th percentile), and max (100th percentile) of the vector `x`. We can use it with summarize like this:

```
pander::pander(heights %>%
  filter(sex == "Female") %>%
  summarize(median_min_max = quantile(height, c(0.5, 0, 1))))
```

median_min_max
64.98
51
79

However, notice that the summaries are returned in a row each. To obtain the results in different columns, we have to define a function that returns a data frame like this:

```
median_min_max <- function(x){
  qs <- quantile(x, c(0.5, 0, 1))
  data.frame(median = qs[1], minimum = qs[2], maximum = qs[3])
}
pander::pander(heights %>%
  filter(sex == "Female") %>%
  summarize(median_min_max(height)))
```

median	minimum	maximum
64.98	51	79

In the next section we learn how useful this approach can be when summarizing by group.

4.7.3 Group then summarize with group_by

A common operation in data exploration is to first split data into groups and then compute summaries for each group. For example, we may want to compute the average and standard deviation for men's and women's heights separately. The `group_by` function helps us do this.

If we type this:

```
pander::pander(head(heights %>% group_by(sex)))
```

sex	height
Male	75
Male	70
Male	68

sex	height
Male	74
Male	61
Female	65

The result does not look very different from heights, except we see Groups: sex [2] when we print the object. Although not immediately obvious from its appearance, this is now a special data frame called a grouped data frame, and dplyr functions, in particular summarize, will behave differently when acting on this object. Conceptually, you can think of this table as many tables, with the same columns but not necessarily the same number of rows, stacked together in one object. When we summarize the data after grouping, this is what happens:

```
pander::pander(heights %>%
  group_by(sex) %>%
  summarize(average = mean(height), standard_deviation = sd(height)))
```

sex	average	standard_deviation
Female	64.94	3.761
Male	69.31	3.611

The summarize function applies the summarization to each group separately.

For another example, let's compute the median, minimum, and maximum murder rate in the four regions of the country using the median_min_max defined above:

```
pander::pander(murders %>%
  group_by(region) %>%
  summarize(median_min_max(rate)))
```

region	median	minimum	maximum
Northeast	1.802	0.3196	3.598
South	3.398	1.457	16.45
North Central	1.971	0.5947	5.36
West	1.292	0.5146	3.63

4.8 pull

The us_murder_rate object defined above represents just one number. Yet we are storing it in a data frame:

```
class(us_murder_rate)
```

```
## [1] "data.frame"
```

since, as most dplyr functions, summarize always returns a data frame.

This might be problematic if we want to use this result with functions that require a numeric value. Here we show a useful trick for accessing values stored in data when using pipes: when a data object is piped that object and its columns can be accessed using the pull function. To understand what we mean take a look at this line of code:

```
us_murder_rate %>% pull(rate)
```

```
## [1] 3.034555
```

This returns the value in the rate column of us_murder_rate making it equivalent to us_murder_rate\$rate.

To get a number from the original data table with one line of code we can type:

```
us_murder_rate <- murders %>%
  summarize(rate = sum(total) / sum(population) * 100000) %>%
  pull(rate)
```

```
us_murder_rate
```

```
## [1] 3.034555
```

which is now a numeric:

```
class(us_murder_rate)
```

```
## [1] "numeric"
```

4.9 Sorting data frames

When examining a dataset, it is often convenient to sort the table by the different columns. We know about the order and sort function, but for ordering entire tables, the dplyr function arrange is useful. For example, here we order the states by population size:

```
pander::pander(murders %>%
  arrange(population) %>%
  head())
```

state	abb	region	population	total	rate
Wyoming	WY	West	563626	5	0.8871
District of Columbia	DC	South	601723	99	16.45
Vermont	VT	Northeast	625741	2	0.3196
North Dakota	ND	North Central	672591	4	0.5947
Alaska	AK	West	710231	19	2.675
South Dakota	SD	North Central	814180	8	0.9826

With arrange we get to decide which column to sort by. To see the states by murder rate, from lowest to highest, we arrange by rate instead:

```
pander::pander(murders %>%
  arrange(rate) %>%
  head())
```

state	abb	region	population	total	rate
Vermont	VT	Northeast	625741	2	0.3196
New Hampshire	NH	Northeast	1316470	5	0.3798
Hawaii	HI	West	1360301	7	0.5146
North Dakota	ND	North Central	672591	4	0.5947
Iowa	IA	North Central	3046355	21	0.6893
Idaho	ID	West	1567582	12	0.7655

Note that the default behavior is to order in ascending order. In dplyr, the function desc transforms a vector so that it is in descending order. To sort the table in descending order, we can type:

```
#murders %>%
  #arrange(desc(rate))
```

4.9.1 Nested sorting

If we are ordering by a column with ties, we can use a second column to break the tie. Similarly, a third column can be used to break ties between first and second and so on. Here we order by region, then within region we order by murder rate:

```
pander::pander(murders %>%
  arrange(region, rate) %>%
  head())
```

state	abb	region	population	total	rate
Vermont	VT	Northeast	625741	2	0.3196
New Hampshire	NH	Northeast	1316470	5	0.3798
Maine	ME	Northeast	1328361	11	0.8281
Rhode Island	RI	Northeast	1052567	16	1.52
Massachusetts	MA	Northeast	6547629	118	1.802
New York	NY	Northeast	19378102	517	2.668

4.9.2 The top n

In the code above, we have used the function head to avoid having the page fill up with the entire dataset. If we want to see a larger proportion, we can use the top_n function. This function takes a data frame as its first argument, the number of rows to show in the second, and the variable to filter by in the third. Here is an example of how to see the top 5 rows:

```
pander::pander(murders %>% top_n(5, rate))
```

state	abb	region	population	total	rate
District of Columbia	DC	South	601723	99	16.45
Louisiana	LA	South	4533372	351	7.743
Maryland	MD	South	5773552	293	5.075
Missouri	MO	North Central	5988927	321	5.36
South Carolina	SC	South	4625364	207	4.475

Note that rows are not sorted by rate, only filtered. If we want to sort, we need to use arrange. Note that if the third argument is left blank, top_n filters by the last column.

4.10 Exercises

For these exercises, we will be using the data from the survey collected by the United States National Center for Health Statistics (NCHS). This center has conducted a series of health and nutrition surveys since the 1960's. Starting in 1999, about 5,000 individuals of all ages have been interviewed every year and they complete the health examination component of the survey. Part of the data is made available via the NHANES package. Once you install the NHANES package, you can load the data like this:

```
library(NHANES)
data(NHANES)
```

The NHANES data has many missing values. The mean and sd functions in R will return NA if any of the entries of the input vector is an NA. Here is an example:

```
library(dslabs)
data(na_example)
mean(na_example)
```

```
## [1] NA
```

```
sd(na_example)
```

```
## [1] NA
```

To ignore the NAs we can use the na.rm argument:

```
mean(na_example, na.rm = TRUE)
```

```
## [1] 2.301754
```

```
sd(na_example, na.rm = TRUE)
```

```
## [1] 1.22338
```

Let's now explore the NHANES data.

1. We will provide some basic facts about blood pressure. First let's select a group to set the standard. We will use 20-to-29-year-old females. AgeDecade is a categorical variable with these ages. Note that the category is coded like " 20-29", with a space in front! What is the average and standard deviation of systolic blood pressure as saved in the BPSysAve variable? Save it to a variable called ref.

Hint: Use filter and summarize and use the na.rm = TRUE argument when computing the average and standard deviation. You can also filter the NA values using filter.

```
NHANES %>%
  filter(AgeDecade == " 20-29") %>%
  summarise(ref = mean(BPSysAve, na.rm = T), ref2 = sd(BPSysAve, na.rm = T))
```

```

## # A tibble: 1 x 2
##      ref    ref2
##      <dbl> <dbl>
## 1    113.   11.7

```

2. Using a pipe, assign the average to a numeric variable ref_avg. Hint: Use the code similar to above and then pull.

```

ref_avg <- NHANES %>%
  filter(AgeDecade == " 20-29") %>%
  summarise(ref = mean(BPSysAve, na.rm = T)) %>%
  pull(ref)

```

3. Now report the min and max values for the same group.

```

NHANES %>%
  filter(AgeDecade == " 20-29") %>%
  summarise(min = min(BPSysAve, na.rm = T), max = max(BPSysAve, na.rm = T))

```

```

## # A tibble: 1 x 2
##      min    max
##      <int> <int>
## 1     84    179

```

4. Compute the average and standard deviation for females, but for each age group separately rather than a selected decade as in question 1. Note that the age groups are defined by AgeDecade. Hint: rather than filtering by age and gender, filter by Gender and then use group_by.

```

NHANES %>%
  filter(Gender == "female") %>%
  group_by(AgeDecade) %>%
  summarise(average = mean(BPSysAve, na.rm = T), standard_deviation = sd(BPSysAve, na.rm = T))

```

```

## # A tibble: 9 x 3
##   AgeDecade average standard_deviation
##   <fct>       <dbl>             <dbl>
## 1 " 0-9"      100.              9.07
## 2 " 10-19"     104.              9.46
## 3 " 20-29"     108.              10.1
## 4 " 30-39"     111.              12.3
## 5 " 40-49"     115.              14.5
## 6 " 50-59"     122.              16.2
## 7 " 60-69"     127.              17.1
## 8 " 70+"       134.              19.8
## 9 <NA>          142.              22.9

```

5. Repeat exercise 4 for males.

```

NHANES %>%
  filter(Gender == "male") %>%
  group_by(AgeDecade) %>%
  summarise(average = mean(BPSysAve, na.rm = T), standard_deviation = sd(BPSysAve, na.rm = T))

```

```

## # A tibble: 9 x 3
##   AgeDecade average standard_deviation
##   <fct>     <dbl>          <dbl>
## 1 "0-9"      97.4           8.32
## 2 "10-19"    110.            11.2
## 3 "20-29"    118.            11.3
## 4 "30-39"    119.            12.3
## 5 "40-49"    121.            14.0
## 6 "50-59"    126.            17.8
## 7 "60-69"    127.            17.5
## 8 "70+"      130.            18.7
## 9 <NA>        136.            23.5

```

6. We can actually combine both summaries for exercises 4 and 5 into one line of code. This is because group_by permits us to group by more than one variable. Obtain one big summary table using group_by(AgeDecade, Gender).

```

NHANES %>%
  group_by(AgeDecade, Gender) %>%
  summarise(average = mean(BPSysAve, na.rm = T), standard_deviation = sd(BPSysAve, na.rm = T))

## `summarise()` has grouped output by 'AgeDecade'. You can override using the '.groups' argument.

## # A tibble: 18 x 4
## # Groups:   AgeDecade [9]
##   AgeDecade Gender average standard_deviation
##   <fct>     <fct>   <dbl>          <dbl>
## 1 "0-9"      female   100.            9.07
## 2 "0-9"      male     97.4           8.32
## 3 "10-19"    female   104.            9.46
## 4 "10-19"    male     110.            11.2
## 5 "20-29"    female   108.            10.1
## 6 "20-29"    male     118.            11.3
## 7 "30-39"    female   111.            12.3
## 8 "30-39"    male     119.            12.3
## 9 "40-49"    female   115.            14.5
## 10 "40-49"   male     121.            14.0
## 11 "50-59"   female   122.            16.2
## 12 "50-59"   male     126.            17.8
## 13 "60-69"   female   127.            17.1
## 14 "60-69"   male     127.            17.5
## 15 "70+"     female   134.            19.8
## 16 "70+"     male     130.            18.7
## 17 <NA>       female   142.            22.9
## 18 <NA>       male     136.            23.5

```

7. For males between the ages of 40-49, compare systolic blood pressure across race as reported in the Race1 variable. Order the resulting table from lowest to highest average systolic blood pressure.

```

NHANES %>%
  filter(Gender == "male" & AgeDecade == "40-49") %>%
  group_by(Race1) %>%
  summarise(average = mean(BPSysAve, na.rm = T)) %>%
  arrange(average)

```

```

## # A tibble: 5 x 2
##   Race1     average
##   <fct>     <dbl>
## 1 White      120.
## 2 Other      120.
## 3 Hispanic   122.
## 4 Mexican    122.
## 5 Black      126.

```

4.11 Tibbles

Tidy data must be stored in data frames. We introduced the data frame in Section 2.4.1 and have been using the murders data frame throughout the book. In Section 4.7.3 we introduced the `group_by` function, which permits stratifying data before computing summary statistics. But where is the group information stored in the data frame?

```
murders %>% group_by(region)
```

```

## # A tibble: 51 x 6
## # Groups:   region [4]
##   state          abb   region   population total   rate
##   <chr>        <chr> <fct>       <dbl> <dbl> <dbl>
## 1 Alabama       AL    South      4779736   135  2.82
## 2 Alaska        AK    West       710231    19   2.68
## 3 Arizona       AZ    West       6392017   232  3.63
## 4 Arkansas      AR    South      2915918   93   3.19
## 5 California    CA    West       37253956  1257 3.37
## 6 Colorado      CO    West       5029196   65   1.29
## 7 Connecticut   CT    Northeast  3574097   97   2.71
## 8 Delaware      DE    South      897934    38   4.23
## 9 District of Columbia DC    South      601723    99  16.5
## 10 Florida      FL   South      19687653  669  3.40
## # ... with 41 more rows

```

Notice that there are no columns with this information. But, if you look closely at the output above, you see the line A tibble followed by dimensions. We can learn the class of the returned object using:

```
murders %>% group_by(region) %>% class()
```

```
## [1] "grouped_df" "tbl_df"      "tbl"        "data.frame"
```

The `tbl`, pronounced `tibble`, is a special kind of data frame. The functions `group_by` and `summarize` always return this type of data frame. The `group_by` function returns a special kind of `tbl`, the `grouped_df`. We will say more about these later. For consistency, the `dplyr` manipulation verbs (`select`, `filter`, `mutate`, and `arrange`) preserve the class of the input: if they receive a regular data frame they return a regular data frame, while if they receive a `tibble` they return a `tibble`. But `tibbles` are the preferred format in the `tidyverse` and as a result `tidyverse` functions that produce a data frame from scratch return a `tibble`. For example, in Chapter 5 we will see that `tidyverse` functions used to import data create `tibbles`.

`Tibbles` are very similar to data frames. In fact, you can think of them as a modern version of data frames. Nonetheless there are three important differences which we describe next.

4.11.1 Tibbles display better

The print method for tibbles is more readable than that of a data frame. To see this, compare the outputs of typing `murders` and the output of `murders` if we convert it to a tibble. We can do this using `as_tibble(murders)`. If using RStudio, output for a tibble adjusts to your window size. To see this, change the width of your R console and notice how more/less columns are shown.

4.11.2 Subsets of tibbles are tibbles

If you subset the columns of a data frame, you may get back an object that is not a data frame, such as a vector or scalar. For example:

```
class(murders[, 4])  
  
## [1] "numeric"
```

is not a data frame. With tibbles this does not happen:

```
class(as_tibble(murders)[, 4])  
  
## [1] "tbl_df"     "tbl"        "data.frame"
```

This is useful in the tidyverse since functions require data frames as input.

With tibbles, if you want to access the vector that defines a column, and not get back a data frame, you need to use the accessor `$`:

```
class(as_tibble(murders)$population)  
  
## [1] "numeric"
```

A related feature is that tibbles will give you a warning if you try to access a column that does not exist. If we accidentally write `Population` instead of `population` this:

```
murders$Population  
  
## NULL
```

returns a `NULL` with no warning, which can make it harder to debug. In contrast, if we try this with a tibble we get an informative warning:

```
as_tibble(murders)$Population  
  
## Warning: Unknown or uninitialized column: 'Population'.  
  
## NULL
```

4.11.3 Tibbles can have complex entries

While data frame columns need to be vectors of numbers, strings, or logical values, tibbles can have more complex objects, such as lists or functions. Also, we can create tibbles with functions:

```
tibble(id = c(1, 2, 3), func = c(mean, median, sd))

## # A tibble: 3 x 2
##       id   func
##   <dbl> <list>
## 1     1 <fn>
## 2     2 <fn>
## 3     3 <fn>
```

4.11.4 Tibbles can be grouped

The function `group_by` returns a special kind of tibble: a grouped tibble. This class stores information that lets you know which rows are in which groups. The tidyverse functions, in particular the `summarize` function, are aware of the group information.

4.11.5 Create a tibble using `tibble` instead of `data.frame`

It is sometimes useful for us to create our own data frames. To create a data frame in the tibble format, you can do this by using the `tibble` function.

```
grades <- tibble(names = c("John", "Juan", "Jean", "Yao"),
                  exam_1 = c(95, 80, 90, 85),
                  exam_2 = c(90, 85, 85, 90))
```

Note that base R (without packages loaded) has a function with a very similar name, `data.frame`, that can be used to create a regular data frame rather than a tibble.

```
grades <- data.frame(names = c("John", "Juan", "Jean", "Yao"),
                      exam_1 = c(95, 80, 90, 85),
                      exam_2 = c(90, 85, 85, 90))
```

To convert a regular data frame to a tibble, you can use the `as_tibble` function.

```
as_tibble(grades) %>% class()

## [1] "tbl_df"     "tbl"        "data.frame"
```

4.12 The dot operator

One of the advantages of using the pipe `%>%` is that we do not have to keep naming new objects as we manipulate the data frame. As a quick reminder, if we want to compute the median murder rate for states in the southern states, instead of typing:

```
tab_1 <- filter(murders, region == "South")
tab_2 <- mutate(tab_1, rate = total / population * 10^5)
rates <- tab_2$rate
median(rates)

## [1] 3.398069
```

We can avoid defining any new intermediate objects by instead typing:

```
filter(murders, region == "South") %>%
  mutate(rate = total / population * 10^5) %>%
  summarize(median = median(rate)) %>%
  pull(median)

## [1] 3.398069
```

We can do this because each of these functions takes a data frame as the first argument. But what if we want to access a component of the data frame. For example, what if the pull function was not available and we wanted to access tab_2\$rate? What data frame name would we use? The answer is the dot operator.

For example to access the rate vector without the pull function we could use

```
rates <- filter(murders, region == "South") %>%
  mutate(rate = total / population * 10^5) %>%
  .$rate
median(rates)

## [1] 3.398069
```

4.13 The purrr package

In Section 3.5 we learned about the sapply function, which permitted us to apply the same function to each element of a vector. We constructed a function and used sapply to compute the sum of the first n integers for several values of n like this:

```
compute_s_n <- function(n){
  x <- 1:n
  sum(x)
}
n <- 1:25
s_n <- sapply(n, compute_s_n)
```

This type of operation, applying the same function or procedure to elements of an object, is quite common in data analysis. The purrr package includes functions similar to sapply but that better interact with other tidyverse functions. The main advantage is that we can better control the output type of functions. In contrast, sapply can return several different object types; for example, we might expect a numeric result from a line of code, but sapply might convert our result to character under some circumstances. purrr functions will never do this: they will return objects of a specified type or return an error if this is not possible.

The first purrr function we will learn is map, which works very similar to sapply but always, without exception, returns a list:

```
library(purrr)
s_n <- map(n, compute_s_n)
class(s_n)

## [1] "list"
```

If we want a numeric vector, we can instead use `map_dbl` which always returns a vector of numeric values.

```
s_n <- map_dbl(n, compute_s_n)
class(s_n)
```

```
## [1] "numeric"
```

This produces the same results as the `sapply` call shown above.

A particularly useful `purrr` function for interacting with the rest of the `tidyverse` is `map_df`, which always returns a tibble data frame. However, the function being called needs to return a vector or a list with names. For this reason, the following code would result in a Argument 1 must have names error:

```
#s_n <- map_df(n, compute_s_n)
```

We need to change the function to make this work:

```
compute_s_n <- function(n){
  x <- 1:n
  tibble(sum = sum(x))
}
s_n <- map_df(n, compute_s_n)
```

The `purrr` package provides much more functionality not covered here. For more details you can consult this online resource.

4.14 Tidyverse conditionals

A typical data analysis will often involve one or more conditional operations. In Section 3.1 we described the `ifelse` function, which we will use extensively in this book. In this section we present two `dplyr` functions that provide further functionality for performing conditional operations.

4.14.1 `case_when`

The `case_when` function is useful for vectorizing conditional statements. It is similar to `ifelse` but can output any number of values, as opposed to just `TRUE` or `FALSE`. Here is an example splitting numbers into negative, positive, and 0:

```
x <- c(-2, -1, 0, 1, 2)
case_when(x < 0 ~ "Negative",
          x > 0 ~ "Positive",
          TRUE ~ "Zero")
## [1] "Negative" "Negative" "Zero"      "Positive"  "Positive"
```

A common use for this function is to define categorical variables based on existing variables. For example, suppose we want to compare the murder rates in four groups of states: New England, West Coast, South, and other. For each state, we need to ask if it is in New England, if it is not we ask if it is in the West Coast, if not we ask if it is in the South, and if not we assign other. Here is how we use `case_when` to do this:

```

murders %>%
  mutate(group = case_when(
    abb %in% c("ME", "NH", "VT", "MA", "RI", "CT") ~ "New England",
    abb %in% c("WA", "OR", "CA") ~ "West Coast",
    region == "South" ~ "South",
    TRUE ~ "Other")) %>%
  group_by(group) %>%
  summarize(rate = sum(total) / sum(population) * 10^5)

```

```

## # A tibble: 4 x 2
##   group      rate
##   <chr>     <dbl>
## 1 New England 1.72
## 2 Other        2.71
## 3 South        3.63
## 4 West Coast   2.90

```

4.14.2 between

A common operation in data analysis is to determine if a value falls inside an interval. We can check this using conditionals. For example, to check if the elements of a vector x are between a and b we can type

```
#x >= a & x <= b
```

However, this can become cumbersome, especially within the tidyverse approach. The between function performs the same operation.

```
#between(x, a, b)
```

4.15 Exercises

- Load the murders dataset. Which of the following is true?

```

data("murders")
head(murders)

```

```

##       state abb region population total
## 1   Alabama  AL   South    4779736   135
## 2    Alaska  AK   West     710231    19
## 3  Arizona  AZ   West    6392017   232
## 4 Arkansas  AR   South    2915918    93
## 5 California CA   West   37253956  1257
## 6 Colorado  CO   West    5029196    65

```

#b. murders is in tidy format and is stored in a data frame.

- Use as_tibble to convert the murders data table into a tibble and save it in an object called murders_tibble.

```
murders_tibble <- as_tibble(murders)
```

3. Use the group_by function to convert murders into a tibble that is grouped by region.

```
murders_tibble %>% group_by(region)
```

```
## # A tibble: 51 x 5
## # Groups:   region [4]
##   state      abb region  population total
##   <chr>     <chr> <fct>    <dbl>   <dbl>
## 1 Alabama    AL   South    4779736   135
## 2 Alaska     AK   West     710231    19
## 3 Arizona    AZ   West     6392017   232
## 4 Arkansas   AR   South    2915918   93
## 5 California CA   West     37253956  1257
## 6 Colorado    CO   West     5029196   65
## 7 Connecticut CT   Northeast 3574097   97
## 8 Delaware    DE   South    897934    38
## 9 District of Columbia DC   South    601723    99
## 10 Florida   FL   South    19687653  669
## # ... with 41 more rows
```

4. Write tidyverse code that is equivalent to this code:

```
#exp(mean(log(murders$population)))
```

Write it using the pipe so that each function is called without arguments. Use the dot operator to access the population. Hint: The code should start with murders %>%.

```
ans <- murders %>% .$population %>%
  log() %>%
  mean() %>%
  exp()

ques <- exp(mean(log(murders$population)))

identical(ans,ques)

## [1] TRUE
```

5. Use the map_df to create a data frame with three columns named n, s_n, and s_n_2. The first column should contain the numbers 1 through 100. The second and third columns should each contain the sum of 1 through n with n the row number.

```
compute_s_n <- function(n){
  x <- 1:n
  tibble(sum = sum(x))
}

data.frame(n = 1:100,
           s_n = map_df(n,compute_s_n),
           s_n_2 = map_df(n,compute_s_n)) %>%
head()
```

```
##   n sum sum.1
## 1 1   1     1
## 2 2   3     3
## 3 3   6     6
## 4 4  10    10
## 5 5  15    15
## 6 6  21    21
```